

```

1 pthread_t thread[MAXTHREADS];
2 struct thread_data data[MAXTHREADS];
3 void pthread_reduce(void) {
4     for (i = 1; k <= nrows - 1; ++k) {
5         for (i = k + 1; i <= nrows; ++i) {
6             data[i] = /* Setup worker. */;
7             pthread_create(&thread[i], NULL, worker,
8                           &data[i]);
9         }
10
11     /* Bug! Should be i <= nrows */
12     for (i = k + 1; i < nrows; ++i)
13         pthread_join(thread[i], NULL);
14 }
15 }

```

Figure 1—pthread gaussian elimination.

```

1 /* Forks a deterministic child. Returns 0 into the
2  * child and 1 into the parent. */
3 int dfork(pid_t childid);
4 /* Merges a child's changes into the parent after
5  * the child issues a dret(). */
6 void djoin(pid_t childid);
7
8 void det_reduce(void) {
9     for (i = 1; k <= nrows - 1; ++k) {
10         for (i = k + 1; i <= nrows; ++i) {
11             data[i] = /* Setup worker. */;
12             if (!dfork(i)) { worker(&data[i]); dret(); }
13         }
14
15     /* Bug! Should be i <= nrows */
16     for (i = k + 1; i < nrows; ++i)
17         djoin(i);
18 }
19 }

```

Figure 2—Deterministic Gaussian elimination.




Figure 3—Data structure relationships associated with object-based reverse mapping. The `struct page` C type encapsulates information about every page frame of physical memory. Two processes map virtual memory to the same page read-only (possibly at different addresses). In order for the kernel to swap a given page to disk, object-based reverse mapping assumes each process maps the page to `vm_area_struct->vm_start + page->index` in virtual memory.




Figure 4—Deterministic speedup for the parallel benchmarks.




Figure 5—Comparing the speedup over  $N = 1$  for the deterministic and pthread versions of the benchmarks. This figure demonstrates the ability of both versions to scale as we add more CPU cores.

Dimension	lu				matmult			
	$N = 1$	$N = 2$	$N = 4$	$N = 8$	$N = 1$	$N = 2$	$N = 4$	$N = 8$
$16 \times 16$	13.1 (41.5%)	45.0 (46.7%)	46.3 (45.8%)	30.9 (31.5%)	9.6 (48.2%)	21.8 (45.0%)	37.8 (45.2%)	16.0 (26.5%)
$32 \times 32$	8.5 (34.0%)	37.3 (45.5%)	45.9 (46.1%)	29.1 (31.1%)	3.3 (37.7%)	13.4 (42.0%)	21.1 (42.2%)	17.5 (24.1%)
$64 \times 64$	2.6 (19.5%)	20.6 (41.6%)	42.1 (44.2%)	32.1 (30.9%)	1.3 (13.4%)	3.8 (26.3%)	7.8 (34.0%)	13.2 (25.8%)
$128 \times 128$	1.4 (2.3%)	6.0 (32.8%)	22.4 (39.0%)	30.8 (31.0%)	1.0 (0.3%)	1.9 (1.7%)	4.5 (13.3%)	6.7 (18.2%)
$256 \times 256$	1.1 (0.5%)	2.1 (11.0%)	7.0 (25.9%)	18.8 (31.1%)	1.0 (0.0%)	1.2 (1.0%)	1.8 (1.6%)	2.3 (5.0%)
$512 \times 512$	1.0 (0.1%)	1.2 (1.4%)	2.3 (9.4%)	5.9 (19.4%)	1.0 (0.0%)	1.0 (0.5%)	1.1 (0.9%)	1.5 (3.0%)
$1024 \times 1024$	1.4 (0.0%)	1.0 (0.3%)	1.2 (1.5%)	1.9 (8.4%)	1.0 (0.0%)	0.9 (0.0%)	1.0 (0.1%)	1.2 (0.2%)
$2048 \times 2048$	1.4 (0.0%)	1.0 (0.0%)	1.0 (0.1%)	1.1 (1.0%)	-	-	-	-

Table 1: Deterministic overhead for *lu* and *matmult*. Overhead is deterministic run time divided by pthread run time. The numbers in parentheses indicate time spent in the kernel performing a virtual memory merge as a percentage of overall runtime.